Outline

2

Outline

3

Regression with ARIMA errors

Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- y_t modeled as function of k explanatory variables
- In regression, we assume that ε_t is white noise.

Regression with ARIMA errors

Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- y_t modeled as function of k explanatory variables
- In regression, we assume that ε_t is white noise.

RegARIMA model

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$

$$\eta_t \sim \text{ARIMA}$$

- Residuals are from ARIMA model.
- Estimate model in one step using MLE
- Select model with lowest AICc value.

US personal consumption and income

```
us_change
```

```
## # A tsibble: 198 x 6 [1Q]
##     Quarter Consumption Income Production Savings Unemployment
##      <qtr>         <dbl>  <dbl>      <dbl>   <dbl>        <dbl>
##  1 1970 Q1        0.619   1.04      -2.45    5.30          0.9
##  2 1970 Q2        0.452   1.23     -0.551    7.79          0.5
##  3 1970 Q3        0.873   1.59     -0.359    7.40          0.5
##  4 1970 Q4       -0.272  -0.240     -2.19    1.17          0.700
##  5 1971 Q1        1.90    1.98       1.91    3.54         -0.100
##  6 1971 Q2        0.915   1.45      0.902    5.87         -0.100
##  7 1971 Q3        0.794   0.521     0.308   -0.406         0.100
##  8 1971 Q4        1.65    1.16       2.29   -1.49          0
##  9 1972 Q1        1.31    0.457      4.15   -4.29         -0.2
## 10 1972 Q2        1.89    1.03       1.89   -4.69         -0.100
## # ... with 188 more rows
```
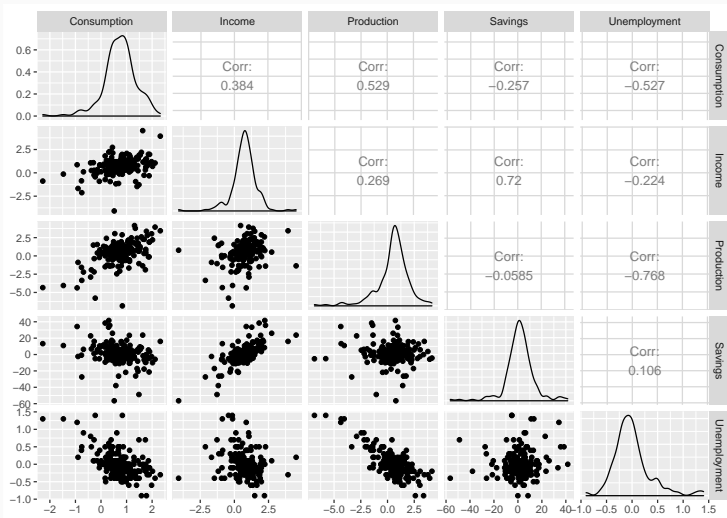
US personal consumption and income

Quarterly changes in US consumption and personal income

US personal consumption and income

```
us_change %>% as_tibble() %>% select(-Quarter) %>% GGally::ggpairs()
```

US personal consumption and income

- No need for transformations or further differencing.
- Increase in income does not necessarily translate into instant increase in consumption (e.g., after the loss of a job, it may take a few months for expenses to be reduced to allow for the new circumstances). We will ignore this for now.

US personal consumption and income

```
fit <- us_change %>%
  model(regarima = ARIMA(Consumption ~ Income + Production + Savings + Unemployment
report(fit)
```

```
## Series: Consumption
## Model: LM w/ ARIMA(0,1,2) errors
##
## Coefficients:
##            ma1     ma2  Income  Production  Savings  Unemployment
##        -1.0882  0.1118  0.7472      0.0370  -0.0531       -0.2096
## s.e.    0.0692  0.0676  0.0403      0.0229   0.0029        0.0986
##
## sigma^2 estimated as 0.09588:  log likelihood=-47.13
## AIC=108.27   AICc=108.86   BIC=131.25
```

# US personal consumption and income

```
fit <- us_change %>%
  model(regarima = ARIMA(Consumption ~ Income + Production + Savings + Unemployment
report(fit)
```

```
## Series: Consumption
## Model: LM w/ ARIMA(0,1,2) errors
##
## Coefficients:
##           ma1     ma2  Income  Production  Savings  Unemployment
##       -1.0882  0.1118  0.7472      0.0370  -0.0531       -0.2096
## s.e.   0.0692  0.0676  0.0403      0.0229   0.0029        0.0986
##
## sigma^2 estimated as 0.09588: log likelihood=-47.13
## AIC=108.27   AICc=108.86   BIC=131.25
```

Write down the equations for the fitted model.

US personal consumption and income

```
gg_tsresiduals(fit)
```

US personal consumption and income

```r
augment(fit) %>%
  features(.resid, ljung_box, dof = 6, lag = 12)
```

```
## # A tibble: 1 x 3
##    .model    lb_stat lb_pvalue
##    <chr>       <dbl>     <dbl>
## 1 regarima     20.0   0.00274
```

US personal consumption and income

```r
us_change_future <- new_data(us_change, 8) %>%
  mutate(Income = tail(us_change$Income,1),
         Production =tail(us_change$Production,1),
         Savings = tail(us_change$Savings,1),
         Unemployment = tail(us_change$Unemployment,1))
forecast(fit, new_data = us_change_future) %>%
  autoplot(us_change) +
  labs(x = "Year", y = "Percentage change",
       title = "Forecasts from regression with ARIMA errors")
```

Forecasting

- To forecast a regression model with ARIMA errors, we need to forecast the regression part of the model and the ARIMA part of the model and combine the results.
- Some predictors are known into the future (e.g., time, dummies).
- Separate forecasting models may be needed for other predictors.
- Forecast intervals ignore the uncertainty in forecasting the predictors.

Daily electricity demand

Model daily electricity demand as a function of temperature using quadratic regression with ARMA errors.

```
vic_elec_daily %>%
  ggplot(aes(x = Temperature, y = Demand, colour = Day_Type)) +
  geom_point() +
  labs(x = "Maximum temperature", y = "Electricity demand (GW)")
```

Daily electricity demand

```
vic_elec_daily %>%
  pivot_longer(c(Demand, Temperature)) %>%
  ggplot(aes(x = Date, y = value)) + geom_line() +
  facet_grid(vars(name), scales = "free_y")
```
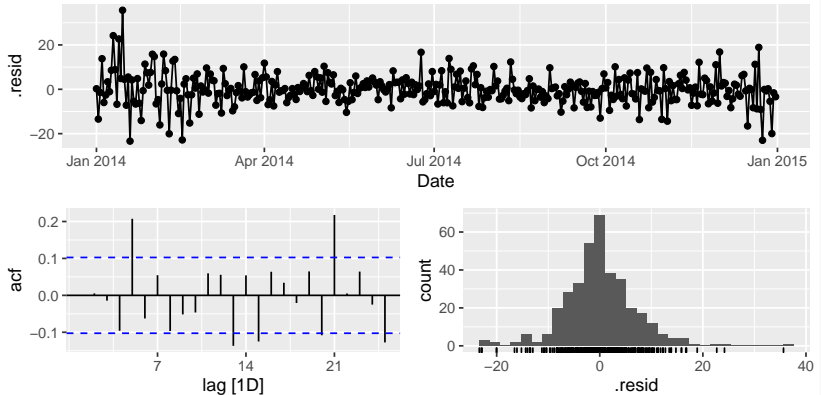
Daily electricity demand

```
fit <- vic_elec_daily %>%
  model(fit = ARIMA(Demand ~ Temperature + I(Temperature^2) +
    (Day_Type == "Weekday")))
report(fit)
```

```
## Series: Demand
## Model: LM w/ ARIMA(2,1,2)(0,0,2)[7] errors
##
## Coefficients:
##          ar1      ar2      ma1     ma2    sma1    sma2  Temperature
##       1.1521  -0.2750  -1.3851  0.4071  0.1589  0.3103      -7.9467
## s.e.  0.6265   0.4812   0.6082  0.5804  0.0591  0.0538       0.4920
##       I(Temperature^2)  Day_Type == "Weekday"TRUE
##                 0.1865                    31.8245
## s.e.            0.0097                     1.0189
##
## sigma^2 estimated as 48.82:  log likelihood=-1220.48
## AIC=2460.96   AICc=2461.58   BIC=2499.93
```

Daily electricity demand

```
augment(fit) %>%
  gg_tsdisplay(.resid, plot_type = "histogram")
```

Daily electricity demand

```
augment(fit) %>%
  features(.resid, ljung_box, dof = 9, lag = 14)
```

```
## # A tibble: 1 x 3
##   .model lb_stat   lb_pvalue
##   <chr>    <dbl>       <dbl>
## 1 fit       38.1 0.000000354
```

Daily electricity demand

```r
# Forecast one day ahead
vic_next_day <- new_data(vic_elec_daily, 1) %>%
  mutate(Temperature = 26, Day_Type = "Holiday")
forecast(fit, vic_next_day)
```

```
## # A fable: 1 x 6 [1D]
## # Key:      .model [1]
##    .model Date       Demand .distribution Temperature Day_Type
##    <chr>  <date>      <dbl> <dist>              <dbl> <chr>
## 1 fit     2015-01-01   161. N(161, 49)             26 Holiday
```

Daily electricity demand

```r
vic_elec_future <- new_data(vic_elec_daily, 14) %>%
  mutate(
    Temperature = 26,
    Holiday = c(TRUE, rep(FALSE, 13)),
    Day_Type = case_when(
      Holiday ~ "Holiday",
      wday(Date) %in% 2:6 ~ "Weekday",
      TRUE ~ "Weekend"
    )
)
```

Daily electricity demand

```
forecast(fit, vic_elec_future) %>%
  autoplot(vic_elec_daily) + ylab("Electricity demand (GW)")
```

Outline

Repeat the daily electricity example, but instead of using a quadratic function of temperature, use a piecewise linear function with the "knot" around 20 degrees Celsius (use predictors `Temperature` & `Temp2`). How can you optimize the choice of knot?

The data can be created as follows.

```
vic_elec_daily <- vic_elec %>%
  filter(year(Time) == 2014) %>%
  index_by(Date = date(Time)) %>%
  summarise(
    Demand = sum(Demand) / 1e3,
    Temperature = max(Temperature),
    Holiday = any(Holiday)
  ) %>%
  mutate(
    Temp2 = I(pmax(Temperature - 20, 0)),
    Day_Type = case_when(
      Holiday ~ "Holiday",
      wday(Date) %in% 2:6 ~ "Weekday",
      TRUE ~ "Weekend"))
```

Outline

24

Dynamic harmonic regression

**Combine Fourier terms with ARIMA errors**

Advantages

- it allows any length seasonality;
- for data with more than one seasonal period, you can include Fourier terms of different frequencies;
- the seasonal pattern is smooth for small values of K (but more wiggly seasonality can be handled by increasing K);
- the short-term dynamics are easily handled with a simple ARMA error.

Disadvantages

- seasonality is assumed to be fixed

Eating-out expenditure

```
aus_cafe <- aus_retail %>%
  filter(
    Industry == "Cafes, restaurants and takeaway food services",
    year(Month) %in% 2004:2018
  ) %>%
  summarise(Turnover = sum(Turnover))
aus_cafe %>% autoplot(Turnover)
```
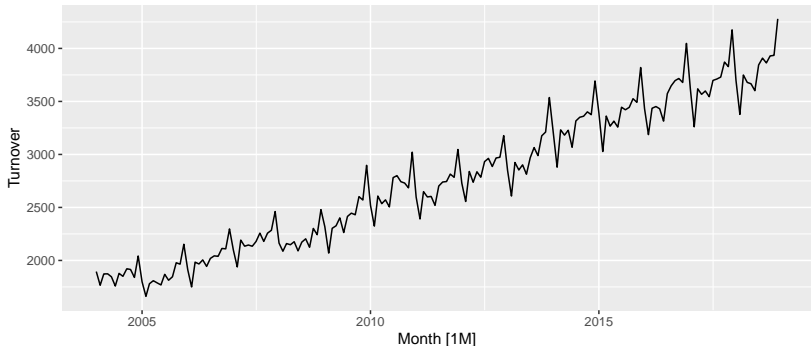
Eating-out expenditure

```
fit <- aus_cafe %>% model(
  K = 1 = ARIMA(log(Turnover) ~ fourier(K = 1) + PDQ(0, 0, 0)),
  K = 2 = ARIMA(log(Turnover) ~ fourier(K = 2) + PDQ(0, 0, 0)),
  K = 3 = ARIMA(log(Turnover) ~ fourier(K = 3) + PDQ(0, 0, 0)),
  K = 4 = ARIMA(log(Turnover) ~ fourier(K = 4) + PDQ(0, 0, 0)),
  K = 5 = ARIMA(log(Turnover) ~ fourier(K = 5) + PDQ(0, 0, 0)),
  K = 6 = ARIMA(log(Turnover) ~ fourier(K = 6) + PDQ(0, 0, 0))
)
glance(fit)
```

.model	sigma2	log_lik	AIC	AICc	BIC
K = 1	0.0017471	317.2353	-616.4707	-615.4056	-587.7842
K = 2	0.0010732	361.8533	-699.7066	-697.8271	-661.4579
K = 3	0.0007609	393.6062	-763.2125	-761.3329	-724.9638
K = 4	0.0005386	426.7839	-821.5678	-818.2098	-770.5697
K = 5	0.0003173	473.7344	-919.4688	-916.9078	-874.8454
K = 6	0.0003163	474.0307	-920.0614	-917.5004	-875.4380

Eating-out expenditure

Eating-out expenditure

Log transformed LM w/ ARIMA(5,1,1) errors, fourier(K = 2)

AICc = −697.8271

Eating-out expenditure

Log transformed LM w/ ARIMA(3,1,1) errors, fourier(K = 3)

AICc = −761.3329

Eating-out expenditure

Eating-out expenditure

Log transformed LM w/ ARIMA(2,1,0) errors, fourier(K = 5)

AICc = −916.9078

Eating-out expenditure

Log transformed LM w/ ARIMA(0,1,1) errors, fourier(K = 6)

AICc = −917.5004

Example: weekly gasoline products

```r
gasoline <- as_tsibble(fpp2::gasoline)
fit <- gasoline %>% model(ARIMA(value ~ fourier(K = 13) + PDQ(0, 0, 0)))
report(fit)
```

```
## Series: value
## Model: LM w/ ARIMA(0,1,1) errors
##
## Coefficients:
##          ma1  fourier(K = 13)C1_52  fourier(K = 13)S1_52  fourier(K = 13)C2_52
##      -0.8934               -0.1121               -0.2300                0.0420
## s.e.  0.0132                0.0123                0.0122                0.0099
##       fourier(K = 13)S2_52  fourier(K = 13)C3_52  fourier(K = 13)S3_52
##                     0.0317                0.0832                0.0346
## s.e.                0.0099                0.0094                0.0094
##       fourier(K = 13)C4_52  fourier(K = 13)S4_52  fourier(K = 13)C5_52
##                     0.0185                0.0398               -0.0315
## s.e.                0.0092                0.0092                0.0091
##       fourier(K = 13)S5_52  fourier(K = 13)C6_52  fourier(K = 13)S6_52
##                     0.0009               -0.0522                0.000
## s.e.                0.0091                0.0090                0.009
##       fourier(K = 13)C7_52  fourier(K = 13)S7_52  fourier(K = 13)C8_52
##                    -0.0173                0.0053                0.0075
## s.e.                0.0090                0.0090                0.0090
##       fourier(K = 13)S8_52  fourier(K = 13)C9_52  fourier(K = 13)S9_52
```

Example: weekly gasoline products

```
forecast(fit, h = "3 years") %>%
  autoplot(gasoline)
```

Outline

Lab Session 19

Repeat Lab Session 18 but using all available data, and handling the annual seasonality using Fourier terms.

Outline

38

Lagged predictors

Sometimes a change in x_t does not affect y_t instantaneously

- y_t = sales, x_t = advertising.
- y_t = stream flow, x_t = rainfall.
- y_t = size of herd, x_t = breeding stock.

Lagged predictors

Sometimes a change in x_t does not affect y_t instantaneously

- y_t = sales, x_t = advertising.
- y_t = stream flow, x_t = rainfall.
- y_t = size of herd, x_t = breeding stock.

- These are dynamic systems with input (x_t) and output (y_t).
- x_t is often a leading indicator.
- There can be multiple predictors.

Lagged predictors

The model include present and past values of predictor: $x_t, x_{t-1}, x_{t-2}, \ldots$.

$$y_t = a + \nu_0 x_t + \nu_1 x_{t-1} + \cdots + \nu_k x_{t-k} + \eta_t$$

where η_t is an ARIMA process.

Lagged predictors

The model include present and past values of predictor: $x_t, x_{t-1}, x_{t-2}, \ldots$.

$$y_t = a + \nu_0 x_t + \nu_1 x_{t-1} + \cdots + \nu_k x_{t-k} + \eta_t$$

where η_t is an ARIMA process.

- x can influence y, but y is not allowed to influence x.

Example: Insurance quotes and TV adverts

```
## # A tsibble: 40 x 3 [1M]
##       Month Quotes TV.advert
##       <mth>  <dbl>    <dbl>
## 1 2002 Jan   13.0     7.21
## 2 2002 Feb   15.4     9.44
## 3 2002 Mar   13.2     7.53
## 4 2002 Apr   13.0     7.21
## 5 2002 May   15.4     9.44
## 6 2002 Jun   11.7     6.42
## 7 2002 Jul   10.1     5.81
## 8 2002 Aug   10.8     6.20
```
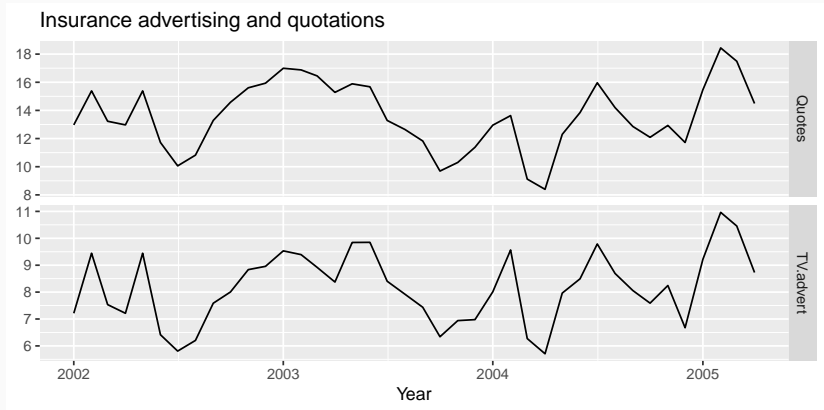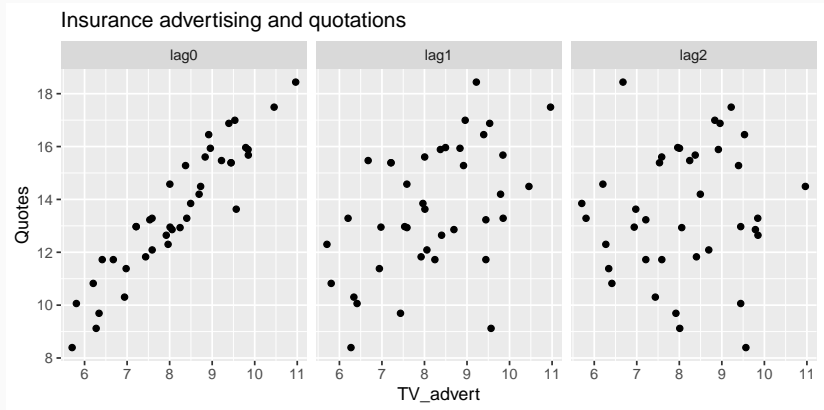
41

Example: Insurance quotes and TV adverts

Insurance advertising and quotations

Example: Insurance quotes and TV adverts

Example: Insurance quotes and TV adverts

```
fit <- insurance %>%
  # Restrict data so models use same fitting period
  mutate(Quotes = c(NA, NA, NA, Quotes[4:40])) %>%
  model(
    ARIMA(Quotes ~ pdq(d = 0) + TV.advert),
    ARIMA(Quotes ~ pdq(d = 0) + TV.advert +
                                lag(TV.advert)),
    ARIMA(Quotes ~ pdq(d = 0) + TV.advert +
                                lag(TV.advert) +
                                lag(TV.advert, 2)),
    ARIMA(Quotes ~ pdq(d = 0) + TV.advert +
                                lag(TV.advert) +
                                lag(TV.advert, 2) +
                                lag(TV.advert, 3))
  )
```

Example: Insurance quotes and TV adverts

```
glance(fit)
```

Lag order	sigma2	log_lik	AIC	AICc	BIC
0	0.2649757	-28.28210	66.56420	68.32890	75.00859
1	0.2094368	-24.04404	58.08809	59.85279	66.53249
2	0.2150429	-24.01627	60.03254	62.57799	70.16581
3	0.2056454	-22.15731	60.31461	64.95977	73.82565

Example: Insurance quotes and TV adverts

```
# Re-fit to all data
fit <- insurance %>%
  model(ARIMA(Quotes ~ TV.advert + lag(TV.advert) + pdq(d = 0)))
report(fit)
```

```
## Series: Quotes
## Model: LM w/ ARIMA(1,0,2) errors
##
## Coefficients:
##          ar1     ma1     ma2  TV.advert  lag(TV.advert)  intercept
##       0.5123  0.9169  0.4591     1.2527          0.1464     2.1554
## s.e.  0.1849  0.2051  0.1895     0.0588          0.0531     0.8595
##
## sigma^2 estimated as 0.2166:  log likelihood=-23.94
## AIC=61.88   AICc=65.38   BIC=73.7
```

Example: Insurance quotes and TV adverts

```
# Re-fit to all data
fit <- insurance %>%
  model(ARIMA(Quotes ~ TV.advert + lag(TV.advert) + pdq(d = 0)))
report(fit)
```

```
## Series: Quotes
## Model: LM w/ ARIMA(1,0,2) errors
##
## Coefficients:
##          ar1     ma1     ma2  TV.advert  lag(TV.advert)  intercept
##       0.5123  0.9169  0.4591     1.2527          0.1464     2.1554
## s.e.  0.1849  0.2051  0.1895     0.0588          0.0531     0.8595
##
## sigma^2 estimated as 0.2166:  log likelihood=-23.94
## AIC=61.88    AICc=65.38    BIC=73.7
```

$$y_t = 2.16 + 1.25x_t + 0.15x_{t-1} + \eta_t,$$
$$\eta_t = 0.512\eta_{t-1} + \varepsilon_t + 0.92\varepsilon_{t-1} + 0.46\varepsilon_{t-2}.$$

Example: Insurance quotes and TV adverts

```
advert_a <- new_data(insurance, 20) %>%
  mutate(TV.advert = 10)
forecast(fit, advert_a) %>% autoplot(insurance)
```

Example: Insurance quotes and TV adverts

```
advert_b <- new_data(insurance, 20) %>%
  mutate(TV.advert = 8)
forecast(fit, advert_b) %>% autoplot(insurance)
```

Example: Insurance quotes and TV adverts

```
advert_c <- new_data(insurance, 20) %>%
  mutate(TV.advert = 6)
forecast(fit, advert_c) %>% autoplot(insurance)
```